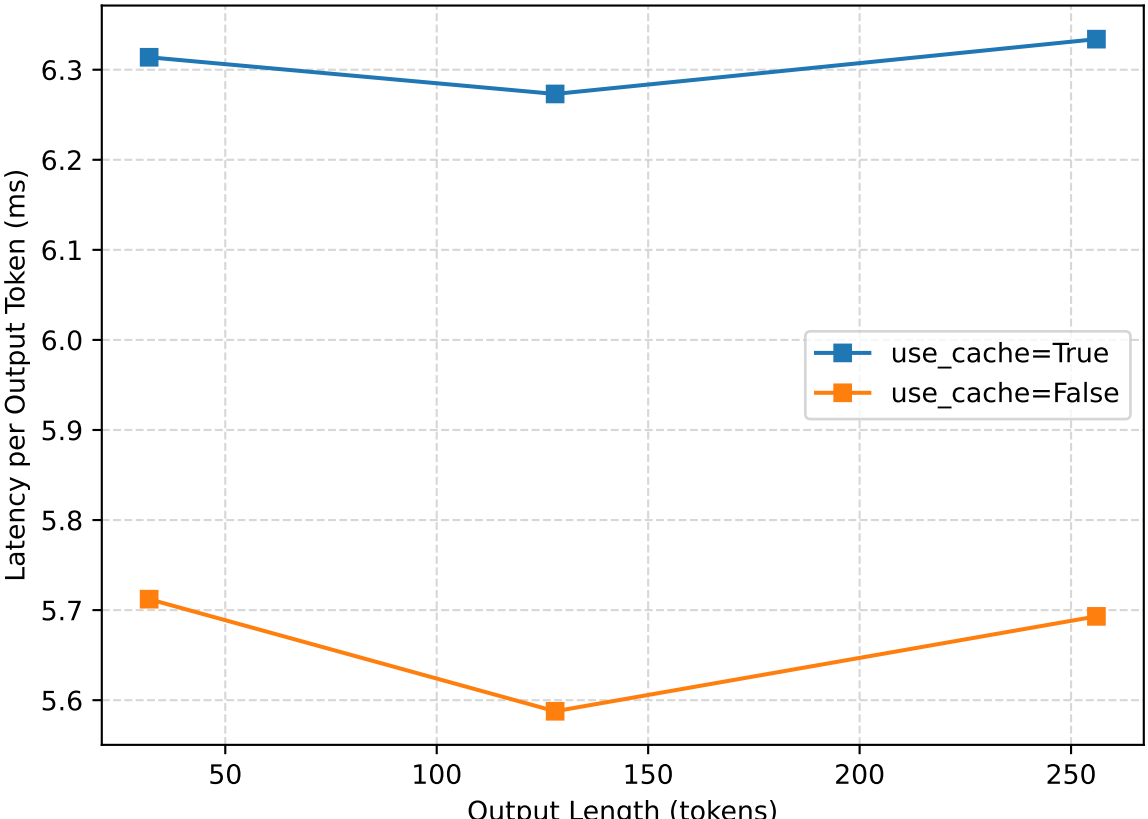
Task 2.2: KV Cache vs. Latency per Output Token
Latency per Output Token (ms)
6.3
6.2
6.1
6.0
5.9
5.8
5.7
5.6
50
100
150
200
250
Output Length (tokens)
use_cache=True
use_cache=False