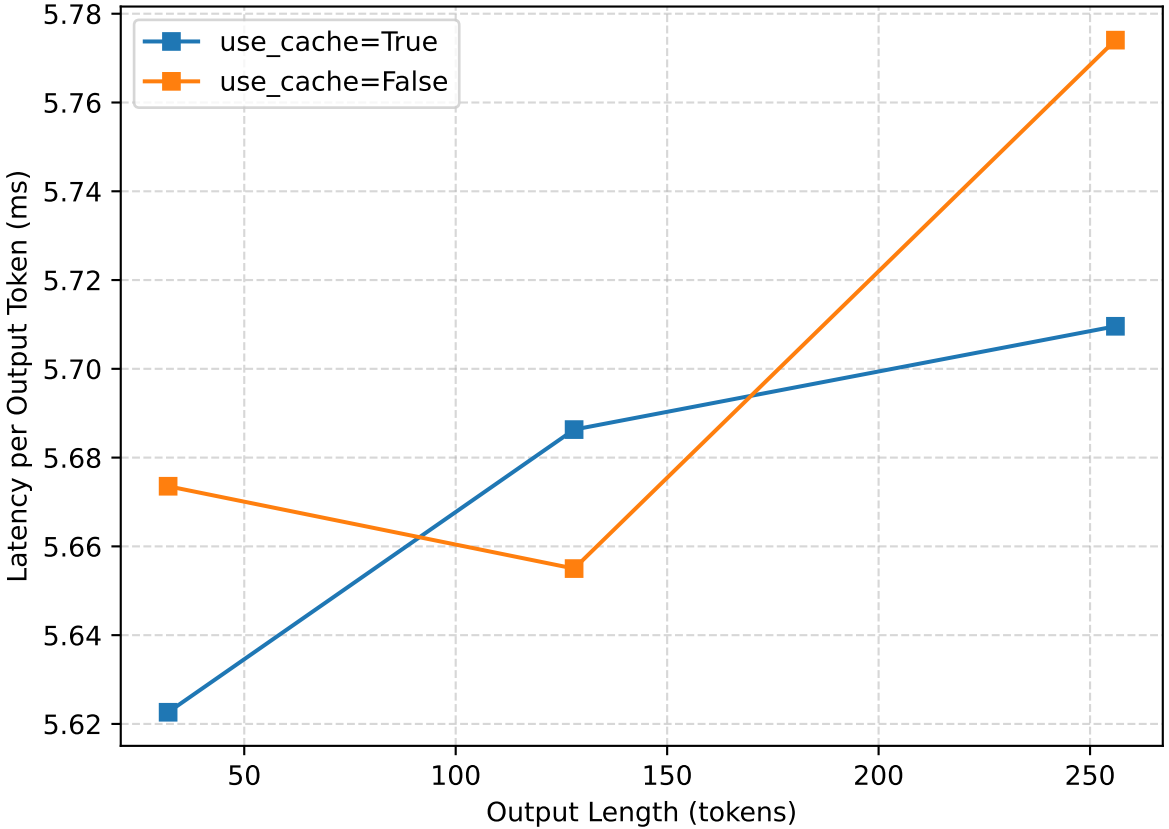
Task 2.2: KV Cache vs. Latency per Output Token
use_cache=True
use_cache=False
Latency per Output Token (ms)
5.78
5.76
5.74
5.72
5.70
5.68
5.66
5.64
5.62
50
100
150
200
250
Output Length (tokens)